

# zebra\_nav

A POMDP navigation environment with a bridge / mine / detour strategy choice, built as a substrate for mechanistic interpretability

## Abstract

ZEBRA\_NAV is a small partially-observable navigation environment in which an agent crosses open procedural terrain to reach a narrow goal door. The direct route is blocked by obstacles the agent can deal with in three distinct ways: **bridge** a lake (place wood on water), **mine** a mountain (turn rock to grass), or **walk around** it. Impassable tree forests line the top and bottom edges so the agent is funnelled through the obstacle-filled middle, where it must repeatedly choose among the three strategies. The point of the design is *not* the navigation itself but the **internal computation it induces**: under partial observation the agent must infer the shape and type of the obstacle ahead (a *belief*) and select a crossing strategy (an *action mode*). We train two agents on the identical task — a recurrent PPO baseline and a pure-JAX DreamerV3 world model — both reaching 100% success, and use them as a testbed for probing and **steering** those belief and strategy representations in activation space. The concrete target is to make an agent that solves the task reliably but **never digs tunnels**, by suppressing the “mine” strategy through an activation-space intervention rather than an action mask.

## 1 Idea and motivation

We want a task that is (i) trivial to solve, so competence is not the bottleneck, yet (ii) forces a small number of *discrete, legible decisions* that depend on a *partially-observed latent* (the obstacle ahead). Such a task is an ideal substrate for mechanistic interpretability: there should be a low-dimensional *belief* subspace (is the obstacle water or rock; how wide is it) and a *strategy* subspace (bridge vs. mine vs. detour) in the agent’s recurrent state, and we can test whether those subspaces are (a) linearly decodable and (b) *causally* steerable. The headline experiment is suppression: keep the agent fully competent while removing one option (tunneling) by editing activations — the recurrent analogue of the “cheese vector” steering result on maze-solving policy networks.

## 2 Environment

**Tiles.** A 9-tile symbolic vocabulary (no decorative classes, so the decoder cannot interpolate phantom tiles):

| id | tile   | role          | behaviour                               |
|----|--------|---------------|-----------------------------------------|
| 0  | GRASS  | open ground   | walkable (cost 1)                       |
| 1  | WATER  | lake          | cross by <i>bridging</i> (place → wood) |
| 2  | ROCK   | mountain      | cross by <i>mining</i> (mine → grass)   |
| 3  | WOOD   | placed bridge | walkable                                |
| 4  | TARGET | goal door     | walkable; stepping on it wins           |
| 5  | OOB    | off-map pad   | egocentric padding only                 |
| 6  | TREE   | forest        | <b>impassable</b> (detour only)         |
| 7  | SAND   | water fringe  | cosmetic, walkable like grass           |
| 8  | DIRT   | rock fringe   | cosmetic, walkable like grass           |

**Observation (POMDP).** An egocentric crop centred on the agent:  $\text{minimap} \in \{0, \dots, 8\}^{V \times V}$  with  $V = 21$  (off-map cells padded with OOB), plus a scalar vector  $\text{scalars} \in \mathbb{R}^5 = [\text{FACING one-hot}(4), \text{step}/\text{max}]$ . The agent never sees the whole map, so the size/type of an obstacle must be *inferred* as it approaches.


**Actions.** `Discrete(6)`: up / down / left / right (a move also sets `FACING`), `PLACE` (bridge the water cell in front: `WATER`  $\rightarrow$  `WOOD`), and `MINE` (mine the rock cell in front: `ROCK`  $\rightarrow$  `GRASS`). Trees can never be removed.

### 3 Task and procedural maps


Maps are  $32 \times 64$ , generated by a domain-warped fractal heightmap thresholded into lakes (`WATER`) and mountains (`ROCK`) with cosmetic `SAND`/ `DIRT` fringes; the left/right edge columns are kept obstacle-free. The **spawn** is the centre of the left edge and the **goal** is a 3-cell central door on the right wall (`goal_half = 1`). Crucially, impassable **tree forests are biased heavily toward the top and bottom rows** (about a 7:1 edge-to-middle ratio), which walls off the trivial “hug the wall to the door” route and forces traversal through the obstacle-filled centre. Figure 1 shows the curated validation set.

zebra\_nav validation maps — natural 32x64

val map 0 (seed 10000)




val map 1 (seed 10001)




val map 2 (seed 10002)



val map 3 (seed 10003)




val map 4 (seed 10004)



val map 5 (seed 10005)




val map 6 (seed 10006)




val map 7 (seed 10007)




val map 8 (seed 10008)




val map 9 (seed 10009)



val map 10 (seed 10010)



val map 11 (seed 10011)



val map 12 (seed 10012)



val map 13 (seed 10013)



### 3.1 Generation algorithm

Each map is produced deterministically from a single integer seed (`generate_zebra_map(orientation="natural", seed)`); the steps are:

1. **Domain-warped fractal heightmap.** A multi-octave OpenSimplex field  $h_0$  (base frequency  $\max(6, \max(H, W)/4.5)$ , five octaves with weights  $1, \frac{1}{2}, \frac{1}{4}, \frac{1}{8}, \frac{1}{16}$ ) is *domain-warped* by two further simplex fields (warp amplitude  $0.10 \max(H, W)$ ):  $h(r, c) = h_0(r + a w_r, c + a w_c)$ . The warp removes grid-aligned, blobby structure and yields organic coastlines.
2. **Overlaid landforms.** 2–4 lakes (subtract a height bump) and 2–4 mountains (add one) are stamped, each a soft logistic *disk* or a multi-lobe *massif*; radii are drawn with a squared (small-biased) distribution so most features are small with a rare large one. A thin wiggly *ridge* (capsule field) is added with probability  $\frac{1}{2}$ .
3. **Quantile thresholding.** The lowest `water_frac` = 0.14 fraction of  $h$  becomes **WATER**, the highest `rock_frac` = 0.14 becomes **ROCK**, the rest **GRASS**. Thresholding by quantile makes obstacle coverage exact regardless of the noise realisation.
4. **Cosmetic fringes.** A sparse one-cell **SAND** speckle is dilated around water and **DIRT** around rock (each kept with prob. 0.30/0.25); both behave exactly like grass but visually mark shore-lines/scree.
5. **Edge bands & endpoints.** The left/right 10 columns are forced to **GRASS** (clean spawn/goal approaches). Spawn =  $(H/2, 0)$ ; the goal is the 3-cell door  $\{H/2-1, H/2, H/2+1\}$  on the right wall.
6. **Edge-biased forests.** About `tree_frac`· $HW$  cells of impassable **TREE** are placed as radius-1–2 patches. Each patch row is sampled with a heavy edge bias — distance-from-nearest-wall  $d = \frac{1}{2}u^3$  with  $u \sim U(0, 1)$ , snapped to the top or bottom edge — so most forest sits in the outer ~15% of rows, leaving the vertical middle open. A breadth-first **reachability guard** removes any patch that would wall the goal off from the spawn.
7. **Caching.** The dataset is generated once and pickled; the pure-JAX env additionally precomputes a Dijkstra cost-to-go field per map so its PBRs potential is a table lookup (and is identical to the PyTorch env by construction).

**Why this geometry.** An earlier variant used the *whole* right wall as the goal; with a free top/bottom border the cheapest route is to hug an edge and sprint right, and both agents collapsed to wall-hugging with no obstacle crossing. Making the goal a narrow central door *and* foresting the edges restores the intended behaviour: the agent must go through the middle and repeatedly choose bridge / mine / detour. The strategy mix is therefore a property of the map geometry, not of the exploration schedule.

## 4 Reward

A move-minimising shaped reward with no per-strategy bias:

$$r_t = \underbrace{-0.01}_{\text{slack}} + \underbrace{\text{shaping}_t}_{\phi = -\text{ctg}} + \underbrace{3.0 \cdot \mathbb{I}[\text{reach door}]}_{\text{terminal}}, \quad \text{shaping}_t = 0.015 (\text{ctg}_{t-1} - \gamma \text{ctg}_t),$$


with  $\gamma = 0.997$ . The potential  $\phi = -\text{ctg}$  uses a **min-action cost-to-go**: entering WATER/ROCK costs 2 (build + move), walkable tiles cost 1, and TREE is impassable;  $\text{ctg}$  is computed once at reset by Dijkstra seeded from the goal cell. With per-build cost 0, crossing and detouring are valued purely by total action count, so which strategy the agent picks at a given obstacle is driven by geometry (is bridging/mining shorter than going around?), exactly the decision we want to study.

## 5 Agents

Both agents train on the *same* task to a shared metric schema. The PyTorch PPO env (`ZebraNavEnv`) and the pure-JAX env (`ZebraNavJaxEnv`) are proven bit-for-bit equivalent (a parity test asserts identical obs / reward / termination over thousands of steps), so the DreamerV3 world model trains on exactly the task the PPO agent does.

### 5.1 Recurrent PPO baseline


A tile-embedding CNN  $\rightarrow$  GRU recurrent policy (Fig. 2). The  $V \times V$  tile crop (9 classes) is embedded (16-d) and augmented with two CoordConv channels, passed through three  $3 \times 3$  ReLU convolutions, flattened, concatenated with the 5 scalars, and projected to a 256-d embedding. A 128-unit GRU carries memory across the episode; **this  $h_t \in \mathbb{R}^{128}$  is the POMDP belief carrier and the primary probing target**. Two linear heads read it: a categorical actor over the 6 actions and a scalar value. There is *no* auxiliary belief head — bridge/mine/detour is driven purely by reward. Entropy 0.045, no annealing (keeps the stochastic policy spread out). Adam ( $\text{lr} = 3 \times 10^{-4}$ ), GAE ( $\gamma=0.997, \lambda=0.95$ ), clip 0.2.



**Figure 2:** PPO+GRU policy. The 5 scalars are the FACING one-hot(4) plus the step fraction. The 6 actions are up/down/left/right/place/mine. The recurrent  $h_t$  is the only place the partially-observed obstacle ahead can be integrated, hence the probing target.

### 5.2 DreamerV3 world model (25M)

The in-tree pure-JAX `purejaxwm` library learns a recurrent state-space model (RSSM) and trains the actor–critic purely *in imagination* (Fig. 3). Because the minimap is symbolic, the encoder/decoder are 4-block MLPs (RMSNorm + SiLU) over the flat  $V \cdot V + 5$  observation vector — no pixel CNN. Heads read the RSSM features  $[h_t, z_t]$ : a decoder  $\hat{o}_t$ , TwoHot reward/continue heads, and the actor/critic optimised on imagined  $\lambda$ -returns with a slow (EMA) critic.



**Figure 3:** High-level DreamerV3. The RSSM rolls forward in latent space; the actor-critic is trained on imagined trajectories ( $\gamma=0.997, \lambda=0.95$ , TwoHot over 255 symlog bins, RetNorm scaling, LaProp + AGC,  $\text{lr} = 4\text{e-}5$ ). The  $h_t$  (3072-d for the 25M preset) + discrete  $z_t$  are the latent we probe.


| agent           | success | episode return | episode length |
|-----------------|---------|----------------|----------------|
| PPO + GRU       | 100%    | $\approx 3.3$  | $\approx 92$   |
| DreamerV3 (25M) | 100%    | $\approx 3.4$  | $\approx 95$   |

## 6 Behaviour: stochastic rollouts

Figure 4 overlays 200 stochastic rollouts per map for each agent. The forests pin the paths into a central corridor; within it the agents bridge (red), mine (yellow), and detour around the larger masses. Figure 5 isolates one clean rollout of each strategy from the PPO agent — the three behaviours we want to find and manipulate internally.

**What counts as a crossing (strict).** We label a trajectory segment as *bridging* or *tunnelling* only when the path **enters an obstacle body and exits the other side**: a maximal run of consecutive path cells whose original terrain is **WATER** (resp. **ROCK**), with at least 2 distinct cells *and* distinct entry/exit land cells. Placing or mining a single block, or poking into a lake and backing out the same side, does *not* count — it is still “avoid”. This is deliberately stricter than counting build/mine events, and it is exactly the segment-labelling rule we will use for the strategy labels in the activation dataset (§8).


atural\_v2\_ppo · 200 stochastic rollouts/map · success 100% · path=darkblue mine=yellow bridge=re amer\_natural\_v2 · 200 stochastic rollouts/map · success 100% · path=darkblue mine=yellow bridge=



(a) PPO + GRU

(b) DreamerV3

**Figure 4:** 200 stochastic rollouts per map. path = dark blue, bridge = red, mine = yellow. Both agents funnel through the forest-walled central corridor and cross obstacles there.




**Figure 5:** One stochastic rollout per strategy (PPO agent), under the strict crossing definition. The crossed lake/mountain body is shaded; the bold segment is the in→out traversal (green ● = entered the body, magenta ★ = exited the far side), the thin blue line is the rest of the path. **Avoid** traverses no obstacle body; **Bridge** crosses a multi-cell lake (18 water cells here); **Tunnel** crosses a mountain (16 rock cells). Over sampled reaching rollouts the mix is roughly avoid:bridge:tunnel  $\approx 1:2:0.8$ , so all three modes are well represented.

## 7 DreamerV3 imagination

Because DreamerV3 learns a world model, we can roll it forward *in imagination* (latent rollouts under the actor, no environment) and decode each imagined latent back to an egocentric view. Figure 6 plays such a dream, decoded to Crafter sprites: warmup frames are real, then the model

imagines moving and crossing obstacles, predicting coherent grass/water/rock/tree terrain.



**Figure 6:** A DREAMERV3 imagined rollout decoded to sprites (green border = real warmup, red = open-loop imagination). The world model imagines coherent terrain — including the edge tree bands — and dreams mining/placing through the middle.

**An encoding caveat (relevant to probing).** The DREAMER obs encodes the minimap as a single ordinal scalar per cell (`tile_id/NUM_TILES`) and the decoder *regresses* it. Reconstructions are therefore blurry and *not* a categorical tile-belief; a cleaner design (one-hot channels + per-cell softmax decoder with cross-entropy) would make tile beliefs a crisp probe target. The PPO agent already uses a proper tile *embedding*.

## 8 Mechanistic-interpretability plan

The agents above are the substrate. The intended programme:

1. **Activation dataset.** Fix one or two maps, sample many stochastic trajectories (logging seed + full metadata), and record the GRU state  $h_t$  (and the pre-recurrent encoder embedding) at every step, labelled with the ground-truth belief (obstacle type/size ahead) and the committed strategy (avoid / bridge / tunnel) at each obstacle.
2. **Probe.** Fit linear probes / difference-of-means directions for the belief and strategy variables; report decoding accuracy by site, and validate with unsupervised structure (PCA / clustering).
3. **Steer.** Add / clamp those directions during fresh rollouts and measure the causal behavioural change (a probe that decodes but does not steer is suspect).
4. **Suppress.** Derive the “mine” direction from the *decision* contrast (mined vs. went-around on the same rock obstacle), apply it conditionally, and verify the agent re-routes via its own bridge/detour policy rather than failing — i.e. a competent agent that **never digs tunnels**, achieved in activation space.

**Reproduce.** Train: `scripts/train_ppo_zebra.py --config models/zebra_nav/natural_centergoal3.yaml` and `scripts/dreamerv3_zebra_nav.py --size 25M`. Visualise: `scripts/zebra_traj_grid.py`, `scripts/zebra.s`, `scripts/viz_dreamer_zebra_imagine.py`. See `docs/zebra_nav.md` and `docs/codebase_map.md`.